

A very brief intro to AI/ML


Jin Guo

Logistics

- Assignment 1 (Due Sep 20)
- Topic Assignment (in slack)

What is Intelligence?

- Draw a sketch to illustrate your understanding of Intelligence.




Does your “intelligence” mean

- Think like people?
- Act like people?
- Think/Behave rationally?
- or

1.7 Einsteins, 2 Maxwells, 5 Faradays and .3 Manhattan Projects?

AI in the Textbook

Study the problem of building agents that can maximize the expected utility given certain constraints.



AI in the News

The New York Times

SCIENCE

Optimism as Artificial Intelligence Pioneers Reunite

Researchers who in the 1960s tackled a field that is still mystifying scientists come together again.

By John Markoff

PRINT EDITION

December 8, 2009, Page D4

TECHNOLOGY

Google's Computing Power Refines Translation Tool

The company's network pushes the limits of translation technology and has become a favored source for millions.

By Miguel Helft

PRINT EDITION Google Can Now Say No to 'Raw Fish Shoes' in 52 Languages | March 9, 2010, Page A1

TECHNOLOGY

Bringing Data Mining Into the Mainstream

A leading data-mining expert explains why the field is moving from a niche into the mainstream of business.

By Steve Lohr

AI in the News

The New York Times


PERSONAL TECH

Amazon Echo, a.k.a. Alexa, Is a Personal Aide in Need of Schooling

The Amazon Echo is an artificially intelligent personal assistant answering to the name Alexa. It's not yet heavy on the "intelligent," but could one day

By Farhad Manjoo

PRINT EDITION Art
Schooling | Jun



TECHNOLOGY

A Facebook Project to Beam Data From Drones Is a Step Closer to Flight

The company said its unmanned aerial vehicle, intended to bring Internet access to remote areas, is ready for tests in the upper atmosphere, most likely in the United States.

By Vindu Goel and Quentin Hardy

PRINT EDITION A Facebook Project to Beam Data From Drones Is a Step Closer to Flight | July 31, 2015, Page B3



AI in the News

The New York Times

OPINION

How Do You Know a Human Wrote This?

Machines are gaining the ability to write, and they are getting terrifyingly good at it.

By Farhad Manjoo

PRINT EDITION


July 30, 2020, Page A22



TECHNOLOGY

How A.I. Steered Doctors Toward a Possible Coronavirus Treatment

Specialists at the London start-up BenevolentAI helped identify the arthritic drug baricitinib, which is now part of a clinical trial.



TECHNOLOGY


A Case for Banning Facial Recognition

A Google research scientist explains why she thinks the police shouldn't use facial recognition software.

By Shira Ovide

PRINT EDITION

June 10, 2020



AI in the Research Frontier

Latent Variable Modelling with Hyperbolic Normalizing Flows

Avishek Joey Bose, Ariella Smofsky, Renjie Liao, Prakash Panangaden, William L. Hamilton

Laplacian Change Point Detection for Dynamic Graphs

Shenyang Huang, Yasmeen Hitti, Guillaume Rabusseau, Reihaneh Rabbany

Algorithmic Improvements for Deep Reinforcement Learning Applied to Interactive Fiction.

Vishal Jain, William Fedus, Hugo Larochelle, Doina Precup, Marc G Bellemare

A Cross-Domain Transferable Neural Coherence Model

Peng Xu, Hamidreza Saghir, Jin Sung Kang, Teng Long, Avishek Joey Bose, Yanshuai Cao, Jackie Chi Kit Cheung

How To Evaluate Your Dialogue System: Probe Tasks as an Alternative for Token-level Evaluation Metrics

Prasanna Parthasarathi, Joelle Pineau, Sarath Chandar

Building reproducible, reusable, and robust machine learning software

Joelle Pineau

Machine Learning

Constructing and/or learning the parameters of a specified model given existing data

- Supervised

Known: Input instances, corresponding labels
Predict: labels using unseen instances

- Unsupervised

Known: Input instances
Recognize input structure, Generate data

- Reinforcement Learning

Supervised Learning



Input					Output
1	37	Yes	No	No	No
2	39	No	Yes	No	No
3	39.2	Yes	No	Yes	Yes
ID	Temperature	Cough	Sore throat	Headache	Flu
Features					

Training the model

$$activation_w(x) = \sum_{i=1} w_i \cdot f_i(x)$$

Linear classifier



Input					Output
1	37	Yes	No	No	No
2	39	No	Yes	No	No
3	39.2	Yes	No	Yes	Yes
ID	Temperature	Cough	Sore throat	Headache	Flu
Features					

Loss Function

- A Quantitative measure of loss when the output label is different from the label assigned by the classifier
- Use to calculate empirical risk of the classifier with respect to the training data

Input					Output
1	37	Yes	No	No	No
2	39	No	Yes	No	No
3	39.2	Yes	No	Yes	Yes
ID	Temperature	Cough	Sore throat	Headache	Flu
Features					

Training the models

$$activation_w(x) = \sum_{i=1} w_i \cdot f_i(x)$$


Linear classifier

$$z = \sum_{i=1} w_i \cdot f_i(x)$$

$$\sigma(z) = 1/(1 + e^{-z})$$

Probabilistic decisions

<http://playground.tensorflow.org/>



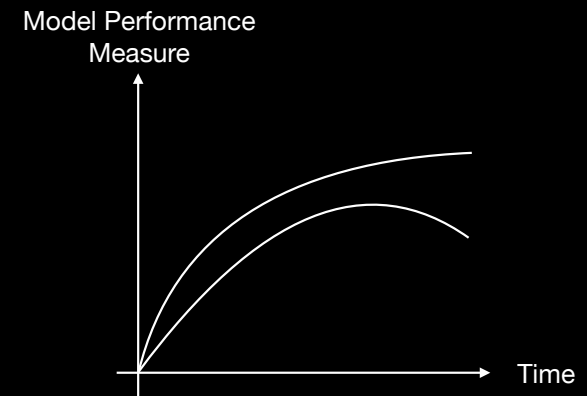

Probabilistic multilayer perceptron

How general is the classifier?

Input					Output
1	37	Yes	No	No	No
2	39	No	Yes	No	No
3	39.2	Yes	No	Yes	Yes
<i>ID</i>	<i>Temperature</i>	<i>Cough</i>	<i>Sore throat</i>	<i>Headache</i>	<i>Flu</i>
<i>Features</i>					


Dataset Split

- Training data
- Validation data (dev, hold-out)
- Test data (consider as final)



Data Preparation

Example: Human Sketch Recognition



Snowman

Ice cream

Giraffe

Grape

Rabbit

Airplane

Tree

Hedgehog


Input



ID

Features


Output



Eitz, M., Hays, J. and Alexa, M., 2012. How do humans sketch objects?. *ACM Transactions on graphics (TOG)*, 31(4), pp.1-10.
<http://cybertron.cg.tu-berlin.de/eitz/projects/classifysketch/>

Data Preparation

Example: Human Sketch Recognition



Input



Features

Eitz, M., Hays, J. and Alexa, M., 2012. How do humans sketch objects?. *ACM Transactions on graphics (TOG)*, 31(4), pp.1-10.
<http://cybertron.cg.tu-berlin.de/eitz/projects/classifysketch/>

Data Preparation

Example: Human Sketch Recognition

